


Julia: A fast & fun dynamic programming language



by Mauro Werder (github: @mauro3, werder@vaw.baug.ethz.ch)

Who am I?

- ▶ Oberassistent in the glaciology group of Daniel Farinotti @ WSL and ETHZ
 - ▶ besides numerical work I also do field and lab experiments

Who am I?

- ▶ Oberassistent in the glaciology group of Daniel Farinotti @ WSL and ETHZ
 - ▶ besides numerical work I also do field and lab experiments
- ▶ Julia user & contributor since 2013 → probably one of the earliest adopters in Switzerland (Julia was first released in 2012)
 - ▶ maintainer of five Julia packages:
 - ▶ Parameters.jl – easier handling of large structs of parameters
 - ▶ UnPack.jl – packing and unpacking of data-structures
 - ▶ SimpleTraits.jl – a simple trait system
 - ▶ WhereTheWaterFlows.jl – a water-routing code
 - ▶ KissMCMC.jl – a Markov chain Monte Carlo sample
- ▶ besides my research I use Julia in teaching at ETHZ:
 - ▶ Physics of Glaciers
 - ▶ Solving PDEs in parallel on GPUs with Julia

<https://eth-vaw-glaciology.github.io/course-101-0250-00/>

What does Julia code look like?

Example, solve Lorenz system of ODEs:

using OrdinaryDiffEq, Plots

function lorenz(x, p, t)

$\sigma = 10$

$\beta = 8/3$

$\rho = 28$

return [$\sigma \cdot (x[2] - x[1])$, $x[1] \cdot (\rho - x[3])$, $x[1] \cdot x[2] - \beta \cdot x[3]$]

end

integrate $dx/dt = \text{lorenz}(t, x)$ numerically from $t=0$ to $t=50$

and IC x_0

$\text{tspan} = (0.0, 50.0)$


$x_0 = [2.0, 0.0, 0.0]$

$\text{sol} = \text{solve}(\text{ODEProblem}(\text{lorenz}, x_0, \text{tspan}), \text{Tsit5}())$

$\text{plot}(\text{sol}, \text{vars}=(1,2,3))$ *# plot Lorenz attractor*

What does Julia code look like?

```
plot(sol, vars=(1,2,3)) # plot Lorenz attractor
```



Julia in brief

Julia 1.0 was released in 2018, current version is 1.8

Features

- ▶ general purpose language with a focus on technical computing
- ▶ dynamic language
 - ▶ interactive development
 - ▶ garbage collection: no manual memory management
- ▶ good performance on par with C & Fortran (through just-ahead-of-time compilation)
 - ▶ No need to vectorize: for-loops are fast
- ▶ multiple dispatch
- ▶ user-defined types are as fast and compact as built-ins
- ▶ Lisp-like macros and other metaprogramming facilities
- ▶ designed for parallelism and distributed computation
- ▶ good inter-op with other languages

Ok, but why?

The two language problem

One language to prototype — one language for production

- ▶ example my institute: GREMS re-write from IDL to C(++)

One language for the users — one language for under-the-hood

- ▶ Numpy (python — C)
- ▶ machine-learning: pytorch, tensorflow

Ok, but why?

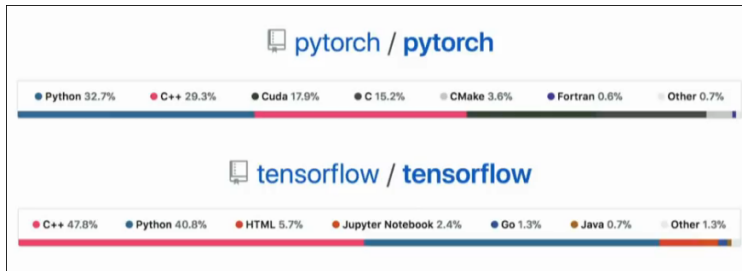
The two language problem

One language to prototype — one language for production

- ▶ example my institute: GREMS re-write from IDL to C(++)

One language for the users — one language for under-the-hood

- ▶ Numpy
- ▶ machine-learning: pytorch, tensorflow



Two language problem

Prototype/interface language:

- ▶ easy to learn and use
- ▶ interactive
- ▶ productive
- ▶ -> **but slow**
- ▶ Examples: Python, Matlab, R, IDL...

Production/fast language:

- ▶ fast
- ▶ -> **but** complicated/verbose/not-interactive/etc
- ▶ Examples: C, C++, Fortran, Java...

Julia solves the two-language problem

Julia is:

- ▶ easy to learn and use
- ▶ interactive
- ▶ productive


and also:

- ▶ fast

This blurs the line between users and developers!

Julia solves the two-language problem

Example:



Flux is a pure Julia deep learning package. This is possible, because Julia runs at native speed both on the CPU and GPU.

The good: easy to learn

- ▶ the basics of Julia are as simple to learn as Python, Matlab, etc
- ▶ however, the language has many powerful/advanced features which can be accessed as the programmer progresses in skills (rich type system, multiple dispatch, meta-programming, access to compiler) → **the language grows with the user**
- ▶ writing code is very close to mathematics (not like the Python/Numpy stuff) → one reason this works so well is *multiple dispatch* (next slide)

The good: nice syntax, cool features

Examples:

Compact function definition and unicode:

$$\phi(t) = \sin(2\pi * t) \ll [1\text{mm}]$$

The good: nice syntax, cool features

Examples:

Compact function definition and unicode:

```
 $\phi(t) = \sin(2\pi * t) \backslash \backslash [1\text{mm}]$ 
```

Vectorized function application:

```
 $\phi.([1,2,3]) .+ [4,5,6]$ 
```

(note, this gets fused into one loop, no temporary arrays are created)

The good: nice syntax, cool features

Examples:

Compact function definition and unicode:

```
φ(t) = sin(2π * t)\\[1mm]
```

Vectorized function application:

```
φ.([1,2,3]) .+ [4,5,6]
```

(note, this gets fused into one loop, no temporary arrays are created)

Parametric types

```
struct Point{N, T}
    coords::NTuple{N,T}
```

```
end
```

```
Point((3,4))    # 2D -> Point{2, Int64}((3, 4))
```

```
Point((3,4,5))  # 3D -> Point{3, Float64}((3.0, 4.0, 5.0))
```

Interlude: multiple dispatch

- ▶ Julia is **not** an object oriented language

Object oriented:

- ▶ methods belong to objects
- ▶ method is selected based on first argument (self)

Multiple dispatch:

- ▶ methods are separate from objects
- ▶ are selected based on all arguments
- ▶ very natural for mathematical programming

Juliacon 2019 presentation on the subject by Stefan Karpinski (co-creator of Julia):
"The Unreasonable Effectiveness of Multiple Dispatch" ([link](#))

Demo-script

```
struct Rock end
struct Paper end
struct Scissors end
## of course structs could have fields as well
# struct Rock
#     color
#     name::String
#     density::Float64
# end

# define multi-method
play(::Rock, ::Paper) = "Paper wins"
play(::Rock, ::Scissors) = "Rock wins"
play(::Scissors, ::Paper) = "Scissors wins"
play(a, b) = play(b, a) # commutative

play(Scissors(), Rock()) # -> "Rock wins"
```

Demo-script (cont.)

```
# Extend-later: with new type
struct Pond end
play(::Rock, ::Pond) = "Pond wins"
play(::Paper, ::Pond) = "Paper wins"
play(::Scissors, ::Pond) = "Pond wins"

play(Scissors(), Pond()) # -> "Pond wins"

# Extend-later: with new function
combine(::Rock, ::Paper) = "Paperweight"
combine(::Paper, ::Scissors) = "Two pieces of papers"
# ...

combine(Rock(), Paper()) # -> "Paperweight"
```

The good: productive

- ▶ Julia is probably as productive or more so than your favourite language.
- ▶ Packages can be easily installed, into project environments if desired.
- ▶ Registered packages are listed on <https://juliahub.com/ui/Packages>
- ▶ The package ecosystem has not reached the level of the Python ecosystem but is growing fast.
There are currently about 7200 registered packages.
→ packages work well together, e.g. combine OrdinaryDiffEq.jl with Unitful.jl and Measurements.jl

The good: inter-op

- ▶ Good interoperability with other languages: `ccall` (C, Fortran), `Cxx.jl`, `PythonCall.jl`, `RCall.jl`, `MATLAB.jl`, etc.

```
ccall((:exp , "libm.so.6"), Cdouble , (Cdouble ,), 1.57) # ->4.806648193775178
```

The good: inter-op

- ▶ Good interoperability with other languages: `ccall` (C, Fortran), `Cxx.jl`, `PythonCall.jl`, `RCall.jl`, `MATLAB.jl`, etc.

```
ccall{(:exp , "libm.so.6"), Cdouble , (Cdouble ,), 1.57} # ->4.806648193775178
```

- ▶ runs on a number of platforms including CPU, GPU, TPU...

on CPU

$f(x) = 3x^2 + 5x + 2$

`C = Array{Float64, 1}([1f0 , 2f0 , 3f0])` *# normal array*

`C .= f.(2 .* C.^2 .+ 6 .* C.^3 .- sqrt.(C))` *# runs on the CPU, note `.`*

on GPU

using CUDA


`B = CuArray{Float64, 1}([1f0 , 2f0 , 3f0])` *# GPU array*

`B .= f.(2 .* B.^2 .+ 6 .* B.^3 .- sqrt.(B))` *# runs on the GPU*

(for CUDA-package to run, you need a Nvidia graphics card)

The good: fast


Julia is as fast as C and Fortan:



Micro-benchmarks from <https://julialang.org/benchmarks/>

The good: fast


Julia is as fast as C and Fortan, but more productive:



Case-study: Celeste.jl

Project to produce an accurate catalogue of 188 million astronomical objects
Ran on Cori supercomputer (@ Berkeley Lab):

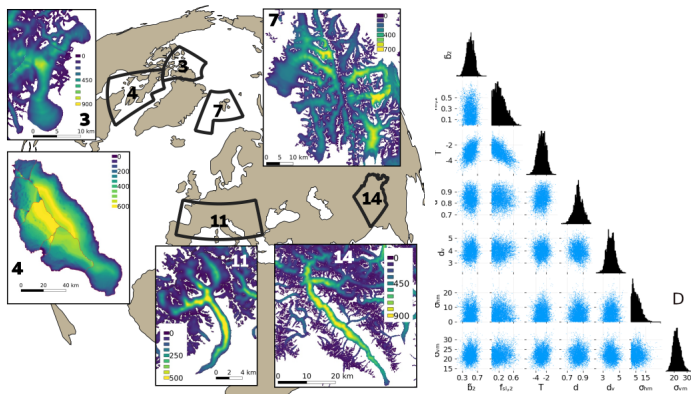
- ▶ 1.54 petaflops using 1.3 million threads on 9,300 Knight Landing (KNL) nodes
→ the first dynamical language to **join the petascale club!**
- ▶ it took 15min to catalogue the objects using Bayesian techniques



Case-study: BITE-model

The Bayesian Ice Thickness Estimation (BITE) model is a project of mine:

- ▶ simple forward model to calculate ice thickness maps
- ▶ using Bayesian techniques (MCMC) to fit it to observations
- ▶ fitted to 30'000 glaciers, calculated 10^8 ice-thickness maps → Julia's performance needed




(Werder et al., 2020) <https://github.com/mauro3/BITEmodel.jl>

Case-study: Julia at CSCS

3D multi physics flow solver in Julia running on Piz Daint by Ludovic Räss (WSL/VAW Glaciology), Samuel Omlin (CSCS) & Yury Podladchikov (Uni Lausanne) (link)

- ▶ original code was a Matlab prototype + CUDA C + MPI production code
- ▶ Julia code running on 5120 NVIDIA Tesla P100 GPUs on the hybrid Cray XC-50 showing nearly perfect scaling



The good: cutting-edge

Whilst the breath of the package ecosystem is not comparable to, say, Python. There are many **cutting-edge** packages:

- ▶ DifferentialEquations.jl, probably the best differential equation solver package around
 - ▶ Flux.jl machine learning/differentiable programming package
 - ▶ JuMP.jl mathematical optimisation package, big in, e.g., operations research
 - ▶ automatic differentiation (AD).
 - ▶ ForwardDiff.jl & ReverseDiff.jl (operator overloading)
 - ▶ Zygote.jl (source transformation)
- above packages can work together to allow *Scientific machine learning*, combining complex physical models with machine learning:
<https://fluxml.ai/blog/2019/03/05/dp-vs-rl.html>
- ▶ etc.

The good: community

Last, the Julia community is very friendly and helpful:

- ▶ Easy to get help on the discourse forum, StackOverflow and on Slack
- ▶ Active developer community on GitHub
- ▶ JuliaCon is as fun and friendly (next one in July fully online and for free)


The bad

- ▶ There are no interface-based abstractions in the language
- ▶ It is not a statically compiled language and thus it does not have the associated safety features.
 - there might be static checkers coming in the future

The ugly


- ▶ Compilation times can be very long. E.g. time to first plot can 20s+ (but fast afterwards)
 - ▶ in general there are improvements from Julia version to version but it is still far from satisfying
 - ▶ apparently Julia 1.9, coming out in the next few weeks, should improve significantly on that again

Conclusions



- ▶ Julia sure is fast and fun
- ▶ solves the two language problem & blurs lines between devs and users

Conclusions



- ▶ Julia sure is fast and fun
- ▶ solves the two language problem & blurs lines between devs and users

Let's get started then!